

Assignment 1

Part 1

a: In your opinion, what were the most important turning points in the history of deep learning?

There have been many important turning points in the history of deep learning. From the development of the first artificial neuron in the 1940s to today, a lot has happened. To me personally, the most important turning points have been the advancements in visual recognition. This is a key moment where deep learning began to show a clear advantage over traditional machine learning.

b: Explain the ADAM optimizer.

The typical optimizer uses some form of gradient descent to find a minimum loss. Problems occur when there are local minima along this path. The algorithm might get stuck and never optimize for minimal loss. Momentum is a way to reduce this problem. The momentum algorithm remembers the direction of previous iterations and pushes in that direction. Another technique that accelerates convergence is by adjusting the learning rate. This is known as the adaptive gradient algorithm, where the learning rate becomes smaller over time as it's divided by the sum of all previous gradients.

The ADAM optimizer combines both of these features, making it one of the fastest optimizers to converge and, therefore, quite convenient. However, studies have shown that standard SGD with momentum can perform better overall.

c: Assume data input is a single 30x40 pixel image. First layer is a convolutional layer with 5 filters, with kernel size 3x2, step size (1,1) and padding='valid'. What are the output dimensions?

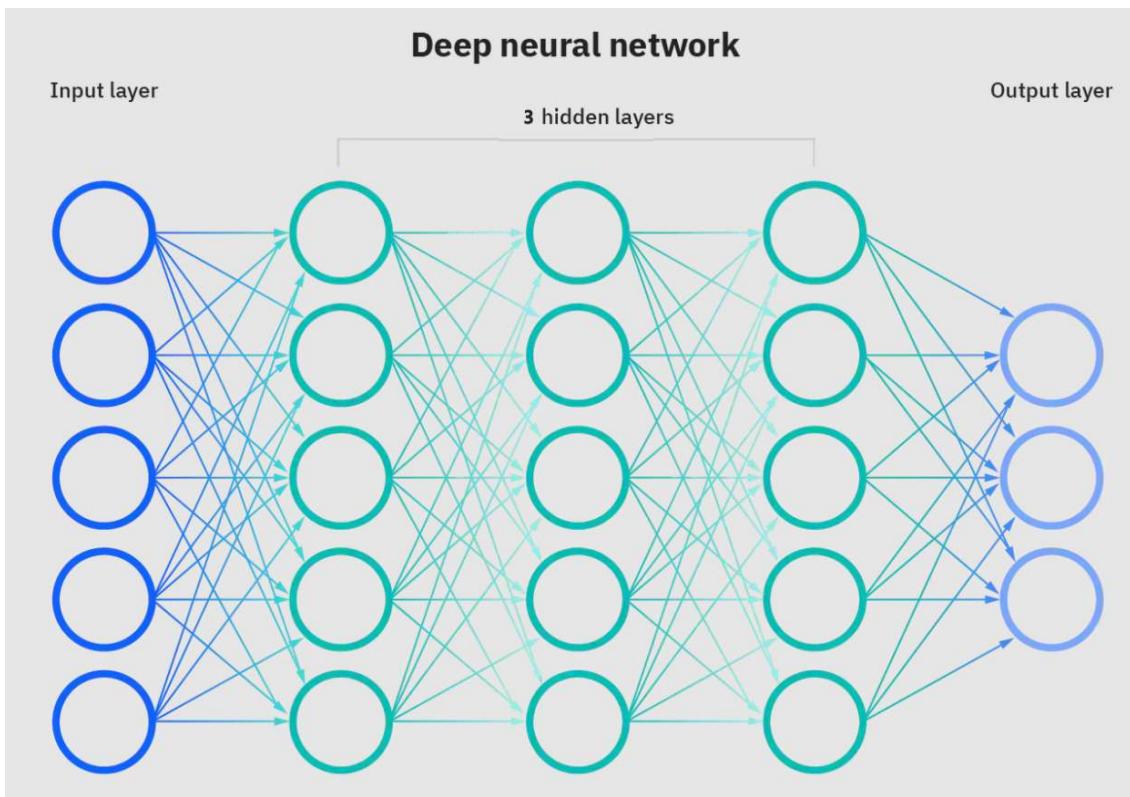
Since the padding is valid, it means that we won't add a padding around the image. The step size is (1,1), so the kernel will only move one step at the time. We have 5 filters, which means that there will be 5 dimensions.

$$\text{Output Height} = \frac{\text{Input Height} - \text{Kernel height}}{\text{Stride Height}} + 1$$

$$\text{Output Width} = \frac{\text{Input Width} - \text{Kernel Width}}{\text{Stride Width}} + 1$$

Following these equations, the output dimensions will be 28x39x5, since we have no padding.

d: Assuming ReLU activations and offsets, and that the last layer is softmax, how many parameters does this network have:



The input layer is 5 and we have 3 hidden layers of size 5 and an output layer of 3. The number of parameters will be:

$$\text{Parameters} = (5 \cdot 5 + 5) + (5 \cdot 5 + 5) + (5 \cdot 5 + 5) + (5 \cdot 3 + 3) = 3 \cdot 30 + 18 = 108$$

e: For a given minibatch, the targets are [1,4, 5, 8] and the network output is [0.1,4.4,0.2,10]. If the loss function is "torch.nn.HuberLoss(reduction='mean', delta=1.0)", what is the loss for this minibatch?

For a batch of size N , the unreduced loss can be described as:

$$L = \{l_1, \dots, l_N\}^T$$

with

$$l_n = \begin{cases} \frac{1}{2}(x_n - y_n)^2, & \text{if } |x_n - y_n| < \delta \\ \delta \left(|x_n - y_n| - \frac{1}{2}\delta \right), & \text{otherwise} \end{cases}$$

If reduction is not none, then:

$$\ell(x, y) = \begin{cases} \text{mean}(L), & \text{if reduction = 'mean'} \\ \text{sum}(L), & \text{if reduction = 'sum'} \end{cases}$$

This will give the result of:

$$L = \left\{ \frac{1}{2}(1 - 0.1)^2, \frac{1}{2}(4 - 4.4)^2, \delta \left(|5 - 0.2| - \frac{1}{2}\delta \right), \delta \left(|8 - 10| - \frac{1}{2}\delta \right) \right\}^T$$

$$L = \{0.405, 0.08, 4.3, 1.5\}^T$$

$$\ell(x, y) = \frac{0.405 + 0.08 + 4.3 + 1.5}{4} = \frac{6.285}{4} = 1.57125$$

The loss of the minibatch will be 1.57125

Part 2: Writing a PyTorch dataset

```
In [1]: import os
import pandas as pd
from PIL import Image
from torch.utils.data import Dataset
from torchvision import transforms
import torch
import matplotlib.pyplot as plt
import numpy as np
from torch import nn

In [2]: class InsectsDataset(Dataset):
    def __init__(self, csv_file, image_folder, root_dir, transform=None):
        """
        Args:
            root_dir (string): Directory with csv and image folder.
            csv_file (string): Path to the csv file with filenames and species.
            image_folder (string): Directory with images.
            transform (callable, optional): Optional transform to be applied on a sample.
        """
        self.root_dir = root_dir
        self.image_folder = image_folder
        csv_path = os.path.join(self.root_dir, csv_file)
        self.insects_df = pd.read_csv(csv_path)
        self.transform = transform

        # Create a dictionary with the species as keys and a unique index as values
        self.species_to_label = {species: idx for idx, species in enumerate(self.insects_df['species'])}
        self.label_to_species = {idx: species for species, idx in self.species_to_label.items()}

    def __len__(self):
        return len(self.insects_df)

    def __getitem__(self, idx):
        # Get the image file path
        img_name = os.path.join(self.root_dir, self.image_folder, self.insects_df.iloc[idx]['image'])
        image = Image.open(img_name)

        # Get the species label
        species = self.insects_df.iloc[idx, 1] # assuming 'species' is the second column
        label = self.species_to_label[species]

        # Apply transformations if needed
        if self.transform:
            image = self.transform(image)

        return image, label

    transform = transforms.Compose([
        transforms.Resize((520, 520)), # Resizing the image to 520x520
        transforms.ToTensor()
    ])

dataset = InsectsDataset(csv_file='insects.csv', image_folder = "Insects", root_dir='data/')

print(dataset.species_to_label)
print(dataset.label_to_species)

{'Andrena fulva': 0, 'Panurgus banksianus': 1, 'Lasioglossum punctatissimum': 2}
{0: 'Andrena fulva', 1: 'Panurgus banksianus', 2: 'Lasioglossum punctatissimum'}
```

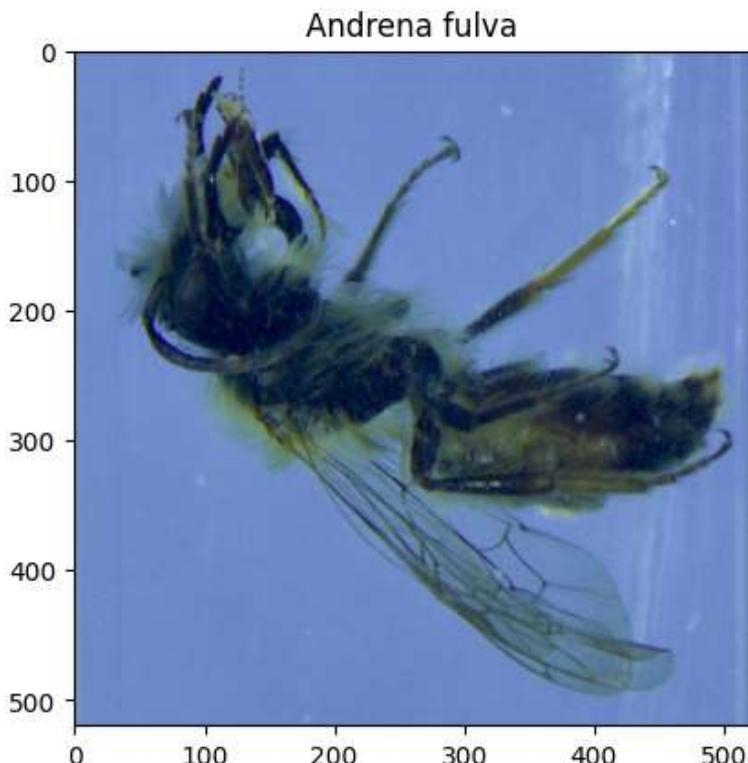
```
In [3]: batch_size = 4

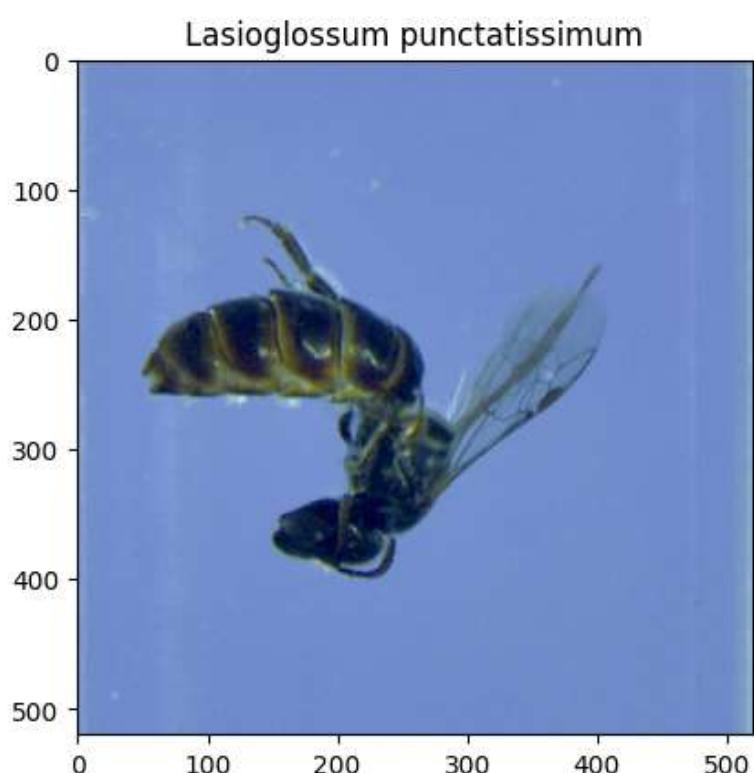
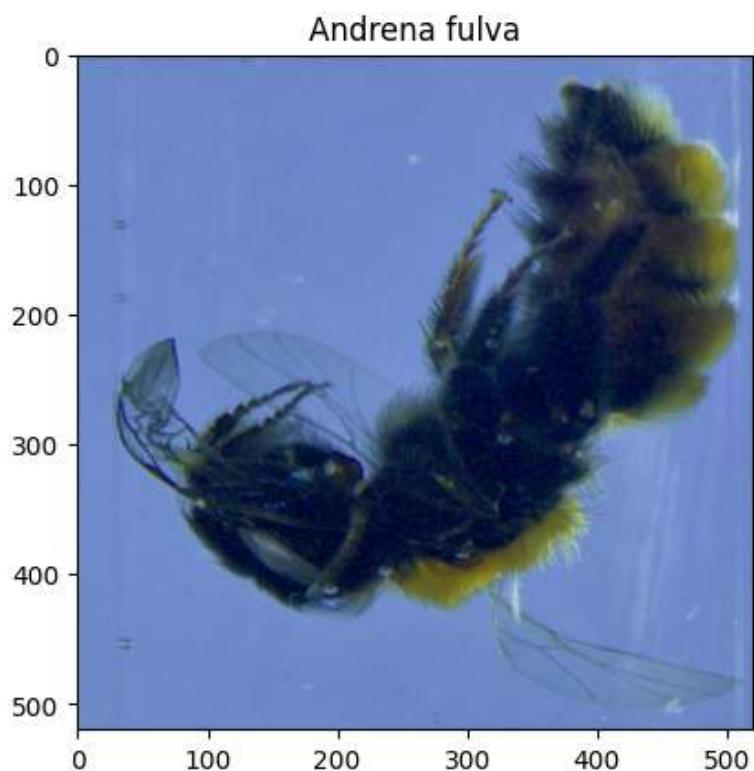
# Set up the dataset.
dataset = dataset

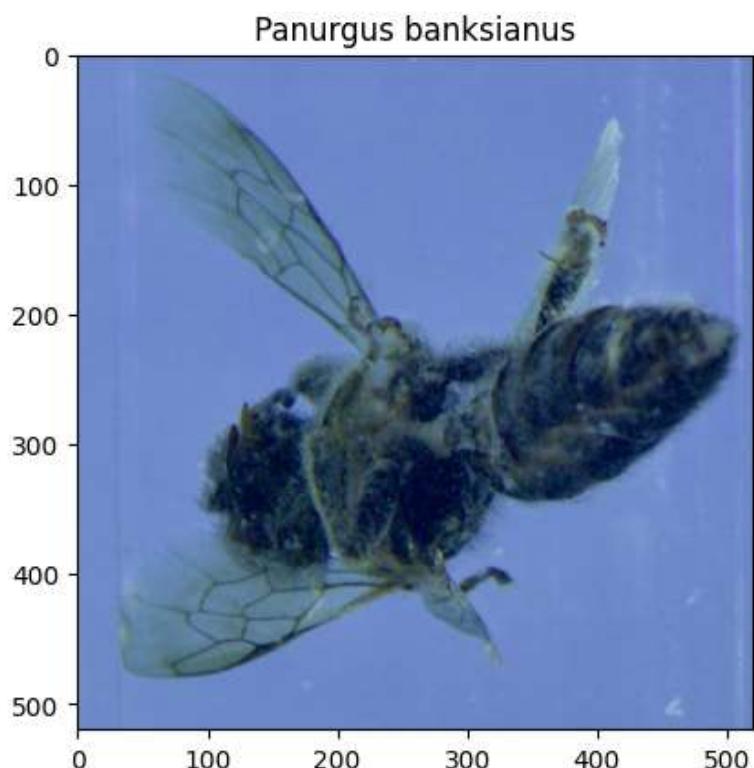
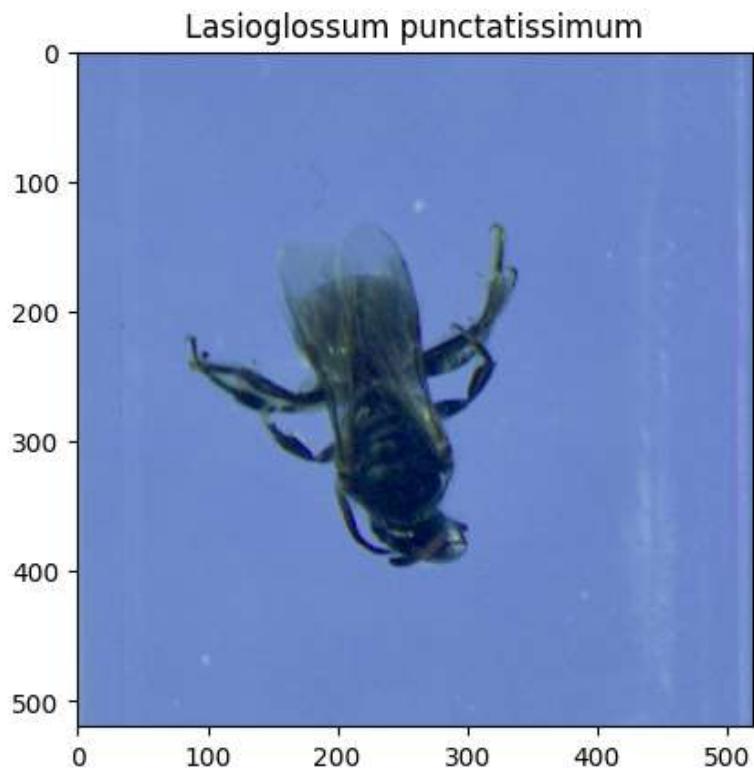
# Set up the dataset.
trainloader = torch.utils.data.DataLoader(dataset,
                                           batch_size=batch_size,
                                           shuffle=True,
                                           num_workers=0)

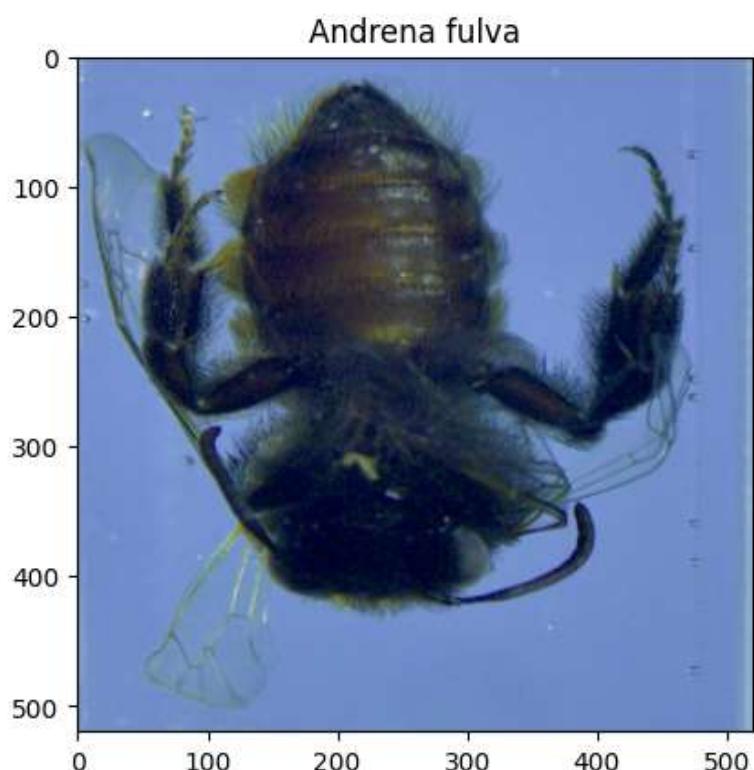
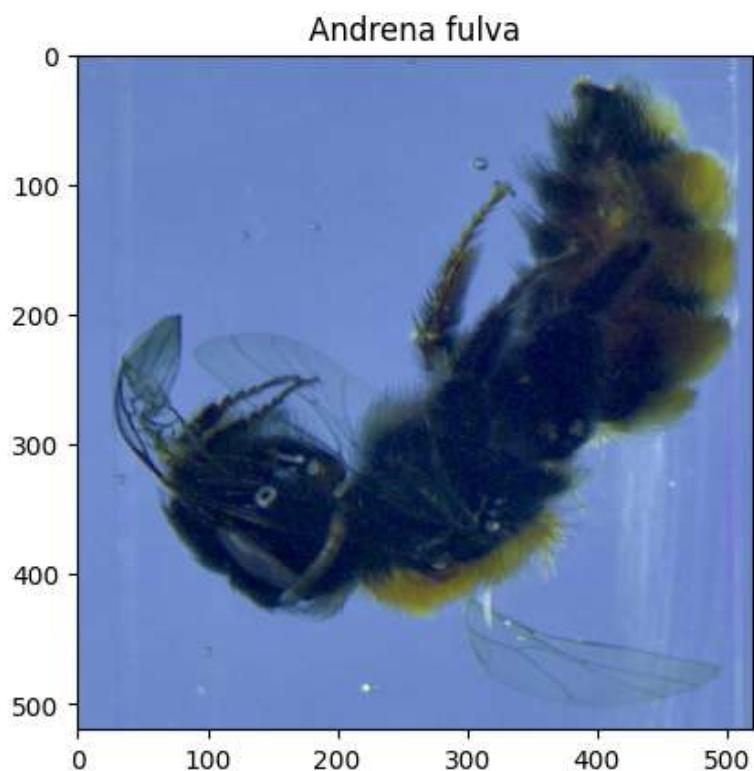
# get some images
dataiter = iter(trainloader)
images, labels = next(dataiter)

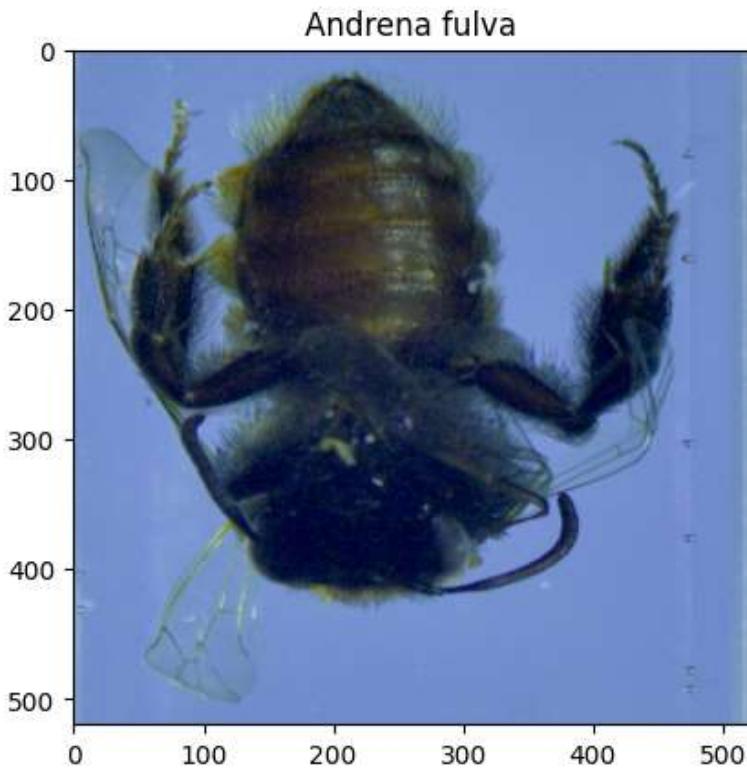
for i in range(2): #Run through 5 batches
    images, labels = next(dataiter)
    for image, label in zip(images,labels): # Run through all samples in a batch
        plt.figure()
        plt.imshow(np.transpose(image.numpy(), (1, 2, 0)))
        plt.title(dataset.label_to_species[label.item()])
```











Part 3:

```
In [4]: class TwoDimensionalData:
    def __init__(self, train_csv, test_csv, root_dir):
        """
        Args:
            root_dir (string): Directory with csv
            csv_file (string): Path to the csv file with filenames and species.
        """
        self.root_dir = root_dir
        train_path = os.path.join(self.root_dir, train_csv)
        test_path = os.path.join(self.root_dir, test_csv)
        self.train_df = pd.read_csv(train_path, sep=";")
        self.test_df = pd.read_csv(test_path, sep=";")

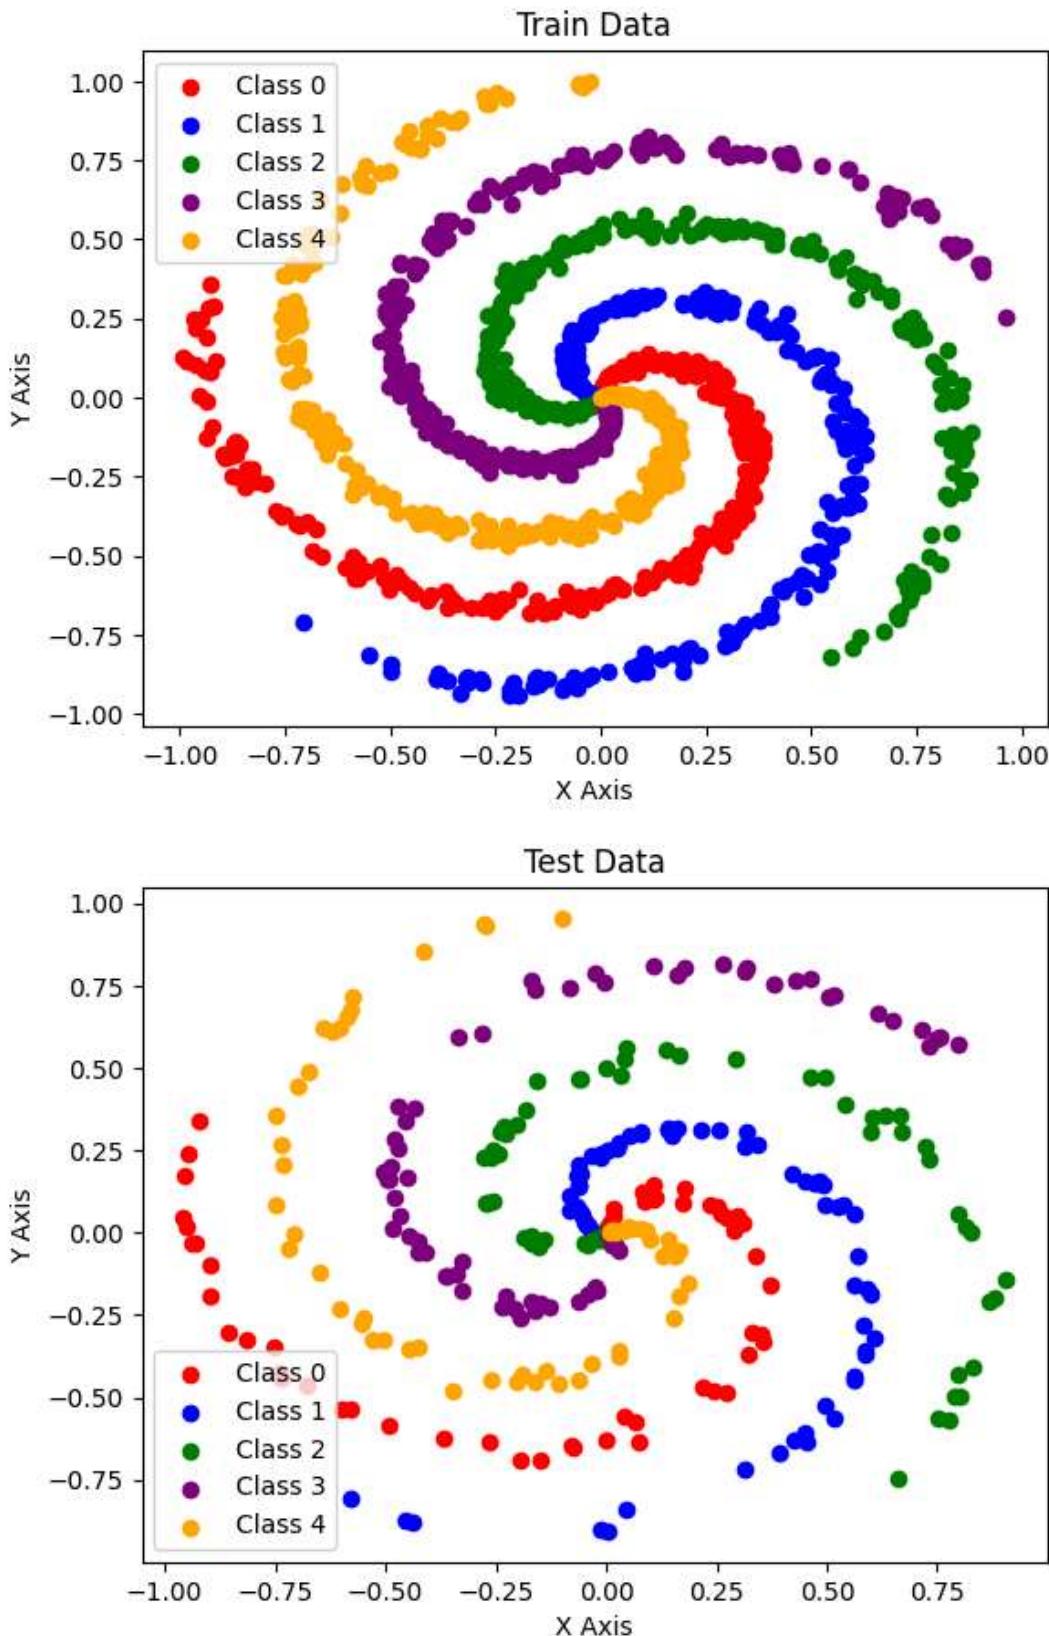
    def plot_data(self, df, title="Data Plot"):
        # Plot x,y with with color based on the Class with matplotlib
        # Use the color_map to assign colors to the classes
        color_map = {0: 'red', 1: 'blue', 2: 'green', 3: 'purple', 4: 'orange'}
        plt.figure()
        for class_value, color in color_map.items():
            plt.scatter(df[df['Class'] == class_value]['X'], df[df['Class'] == class_value]['Y'])
        plt.title(title)
        plt.xlabel('X Axis')
        plt.ylabel('Y Axis')
        plt.legend()
        plt.show()

    def get_data_to_tensor(self, df):
        # Convert the data to a tensor
        data = torch.tensor(df[['X', 'Y']].values, dtype=torch.float32)
        labels = torch.tensor(df['Class'].values, dtype=torch.long)
        return data, labels
twoD = TwoDimensionalData('train_data.csv', 'test_data.csv', 'data/')
```

a: describe & visualize the data

The data looks like a Spiral galaxy, with five arms. Each of the arms have a different class. The data is plotted below.

```
In [5]: twoD.plot_data(twoD.train_df, "Train Data")
twoD.plot_data(twoD.test_df, title="Test Data")
```



b: design a neural network using pytorch to correctly assign labels

```
In [6]: class NNclassifier(nn.Module):
    def __init__(self, input_size, num_classes):
```

```

super(NNclassifier, self).__init__()
# Linear Layers
self.fc1 = nn.Linear(input_size, 256)
self.fc2 = nn.Linear(256, 64)
self.fc3 = nn.Linear(64, num_classes)

# Activation function
self.relu = nn.ReLU()

def forward(self, x):
    out = self.fc1(x)
    out = self.relu(out)
    out = self.fc2(out)
    out = self.relu(out)
    out = self.fc3(out)
    return out

# Hyperparameters
input_size = 2
num_classes = 5
learning_rate = 0.001

# Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Load the data
dataset = TwoDimensionalData('train_data.csv', 'test_data.csv', 'data/')
train_data, train_labels = dataset.get_data_to_tensor(dataset.train_df)
test_data, test_labels = dataset.get_data_to_tensor(dataset.test_df)

train_dataset = torch.utils.data.TensorDataset(train_data, train_labels)
test_dataset = torch.utils.data.TensorDataset(test_data, test_labels)

train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_dataloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

# Initialize the neural network
model = NNclassifier(input_size, num_classes).to(device)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

class trainNN:
    train_losses = []
    test_losses = []
    train_accuracies = []
    test_accuracies = []
    accuracy = []
    convergence = False

    def __init__(self, model, train_dataloader, test_dataloader, criterion, optimizer, num_epochs):
        self.model = model
        self.train_dataloader = train_dataloader
        self.test_dataloader = test_dataloader
        self.criterion = criterion
        self.optimizer = optimizer
        self.num_epochs = num_epochs

    def _train(self, dataloader, model, loss_fn, optimizer):
        size = len(dataloader.dataset)
        model.train()
        total_loss, correct = 0, 0 # Track the total loss for the epoch

        for _, (X, y) in enumerate(dataloader):
            X, y = X.to(device), y.to(device)

```

```

        # Compute prediction error
        pred = model(X)
        loss = loss_fn(pred, y)
        total_loss += loss.item() # Accumulate Loss

        # Backpropagation
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

        # Calculate accuracy
        correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    avg_loss = total_loss / len(dataloader) # Calculate average Loss for the epoch
    avg_accuracy = correct / size

    self.train_losses.append(avg_loss) # Store the average Loss for the epoch
    self.train_accuracies.append(avg_accuracy) # Store the average accuracy for the epoch
    print(f"Train loss: {avg_loss:>7f}, Accuracy: {(100*avg_accuracy):>0.1f}%")

def _test(self, dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    model.eval()
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()
    test_loss /= num_batches
    correct /= size
    self.test_losses.append(test_loss) # Store the test Loss for the epoch
    self.test_accuracies.append(correct)

    print(f"Test Error: \nAvg loss: {test_loss:>8f}, Accuracy: {(100*correct):>0.1f}%")
    return float(correct*100)

def _plot_training_result(self):
    plt.subplot(2,1,1)
    plt.plot(range(1, self.num_epochs+1), self.train_losses, label="Train Loss")
    plt.plot(range(1, self.num_epochs+1), self.test_losses, label="Test Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.title("Training and Testing Loss over Epochs")
    plt.grid()
    plt.legend()
    plt.show()

    plt.subplot(2,1,2)
    plt.plot(range(1, self.num_epochs+1), self.train_accuracies, label="Train Accuracy")
    plt.plot(range(1, self.num_epochs+1), self.test_accuracies, label="Test Accuracy")
    plt.xlabel("Epochs")
    plt.ylabel("Accuracy")
    plt.title("Training and Testing Accuracy over Epochs")
    plt.grid()
    plt.legend()
    plt.show()

def train_model(self):
    accuracy = []
    convergence = False
    for t in range(self.num_epochs):
        print(f"Epoch {t+1}\n-----")

```

```
        self._train(self.train_dataloader, self.model, self.criterion, self.optimizer)
        test_acc = self._test(self.test_dataloader, self.model, self.criterion)
        accuracy.append(test_acc)
        # Check for convergence, meaning the test loss has not decreased for 30 epochs
        if t > 31 and self.test_losses[-1] >= np.mean(self.test_losses[-31:-1]) and not
           convergence:
            convergence = True
        if convergence and accuracy[-1] >= np.max(accuracy):
            self.num_epochs = t + 1
            break
    print("Done!")
    self._plot_training_result()
```

```
In [7]: test = trainNN(model, train_dataloader, test_dataloader, criterion, optimizer, 1000)
test.train_model()
```

Epoch 1

Train loss: 1.412095, Accuracy: 30.2%
Test Error:
Avg loss: 1.253044, Accuracy: 42.1%

Epoch 2

Train loss: 1.070897, Accuracy: 55.3%
Test Error:
Avg loss: 0.870811, Accuracy: 65.6%

Epoch 3

Train loss: 0.756207, Accuracy: 72.8%
Test Error:
Avg loss: 0.630442, Accuracy: 76.9%

Epoch 4

Train loss: 0.555226, Accuracy: 79.5%
Test Error:
Avg loss: 0.467917, Accuracy: 85.3%

Epoch 5

Train loss: 0.428471, Accuracy: 85.4%
Test Error:
Avg loss: 0.388163, Accuracy: 87.6%

Epoch 6

Train loss: 0.349995, Accuracy: 86.8%
Test Error:
Avg loss: 0.302842, Accuracy: 89.0%

Epoch 7

Train loss: 0.277726, Accuracy: 90.2%
Test Error:
Avg loss: 0.281831, Accuracy: 91.3%

Epoch 8

Train loss: 0.233534, Accuracy: 92.6%
Test Error:
Avg loss: 0.211579, Accuracy: 91.3%

Epoch 9

Train loss: 0.210135, Accuracy: 92.1%
Test Error:
Avg loss: 0.202949, Accuracy: 91.6%

Epoch 10

Train loss: 0.183296, Accuracy: 93.7%
Test Error:
Avg loss: 0.184218, Accuracy: 91.3%

Epoch 11

Train loss: 0.160243, Accuracy: 94.5%
Test Error:
Avg loss: 0.167812, Accuracy: 94.3%

Epoch 12

Train loss: 0.147170, Accuracy: 95.6%
Test Error:
Avg loss: 0.150958, Accuracy: 94.3%

Epoch 13

Train loss: 0.139857, Accuracy: 94.3%
Test Error:
Avg loss: 0.132336, Accuracy: 96.3%

Epoch 14

Train loss: 0.116900, Accuracy: 95.7%
Test Error:
Avg loss: 0.126748, Accuracy: 95.3%

Epoch 15

Train loss: 0.110410, Accuracy: 96.2%
Test Error:
Avg loss: 0.118313, Accuracy: 95.7%

Epoch 16

Train loss: 0.100038, Accuracy: 96.7%
Test Error:
Avg loss: 0.131014, Accuracy: 95.0%

Epoch 17

Train loss: 0.100529, Accuracy: 96.5%
Test Error:
Avg loss: 0.108257, Accuracy: 95.7%

Epoch 18

Train loss: 0.088449, Accuracy: 96.8%
Test Error:
Avg loss: 0.093085, Accuracy: 97.0%

Epoch 19

Train loss: 0.080294, Accuracy: 97.8%
Test Error:
Avg loss: 0.100391, Accuracy: 96.7%

Epoch 20

Train loss: 0.084843, Accuracy: 97.4%
Test Error:
Avg loss: 0.088513, Accuracy: 96.3%

Epoch 21

Train loss: 0.072644, Accuracy: 97.8%
Test Error:
Avg loss: 0.085064, Accuracy: 97.0%

Epoch 22

Train loss: 0.073089, Accuracy: 97.2%
Test Error:
Avg loss: 0.071873, Accuracy: 98.3%

Epoch 23

Train loss: 0.065152, Accuracy: 97.8%
Test Error:
Avg loss: 0.075005, Accuracy: 97.3%

Epoch 24

Train loss: 0.058062, Accuracy: 98.4%
Test Error:
Avg loss: 0.072698, Accuracy: 97.7%

Epoch 25

Train loss: 0.068026, Accuracy: 97.1%
Test Error:
Avg loss: 0.111721, Accuracy: 95.3%

Epoch 26

Train loss: 0.059478, Accuracy: 98.1%
Test Error:
Avg loss: 0.059249, Accuracy: 99.0%

Epoch 27

Train loss: 0.055335, Accuracy: 98.1%
Test Error:
Avg loss: 0.116745, Accuracy: 96.0%

Epoch 28

Train loss: 0.056902, Accuracy: 98.2%
Test Error:
Avg loss: 0.072702, Accuracy: 97.0%

Epoch 29

Train loss: 0.055891, Accuracy: 98.2%
Test Error:
Avg loss: 0.069560, Accuracy: 97.0%

Epoch 30

Train loss: 0.052457, Accuracy: 98.3%
Test Error:
Avg loss: 0.050184, Accuracy: 98.7%

Epoch 31

Train loss: 0.046009, Accuracy: 98.4%
Test Error:
Avg loss: 0.048548, Accuracy: 99.0%

Epoch 32

Train loss: 0.056049, Accuracy: 98.1%
Test Error:
Avg loss: 0.053807, Accuracy: 98.0%

Epoch 33

Train loss: 0.041201, Accuracy: 98.7%
Test Error:
Avg loss: 0.082201, Accuracy: 96.7%

Epoch 34

Train loss: 0.048475, Accuracy: 98.7%
Test Error:
Avg loss: 0.059776, Accuracy: 97.7%

Epoch 35

Train loss: 0.047802, Accuracy: 98.1%
Test Error:
Avg loss: 0.049374, Accuracy: 98.0%

Epoch 36

Train loss: 0.045964, Accuracy: 98.1%
Test Error:
Avg loss: 0.057413, Accuracy: 97.7%

Epoch 37

Train loss: 0.038544, Accuracy: 98.8%
Test Error:
Avg loss: 0.042291, Accuracy: 98.7%

Epoch 38

Train loss: 0.037254, Accuracy: 99.0%
Test Error:
Avg loss: 0.077234, Accuracy: 97.0%

Epoch 39

Train loss: 0.036103, Accuracy: 98.9%
Test Error:
Avg loss: 0.046188, Accuracy: 98.7%

Epoch 40

Train loss: 0.039386, Accuracy: 98.7%
Test Error:
Avg loss: 0.061046, Accuracy: 97.3%

Epoch 41

Train loss: 0.036142, Accuracy: 98.8%
Test Error:
Avg loss: 0.065662, Accuracy: 97.7%

Epoch 42

Train loss: 0.068177, Accuracy: 97.7%
Test Error:
Avg loss: 0.042533, Accuracy: 98.3%

Epoch 43

Train loss: 0.036268, Accuracy: 98.8%
Test Error:
Avg loss: 0.050800, Accuracy: 97.7%

Epoch 44

Train loss: 0.034831, Accuracy: 98.9%
Test Error:
Avg loss: 0.064737, Accuracy: 97.3%

Epoch 45

Train loss: 0.051191, Accuracy: 98.4%
Test Error:
Avg loss: 0.038016, Accuracy: 99.0%

Epoch 46

Train loss: 0.031838, Accuracy: 99.0%
Test Error:
Avg loss: 0.042148, Accuracy: 98.3%

Epoch 47

Train loss: 0.030869, Accuracy: 99.2%
Test Error:
Avg loss: 0.057429, Accuracy: 98.3%

Epoch 48

Train loss: 0.031805, Accuracy: 98.6%
Test Error:
Avg loss: 0.049654, Accuracy: 98.3%

Epoch 49

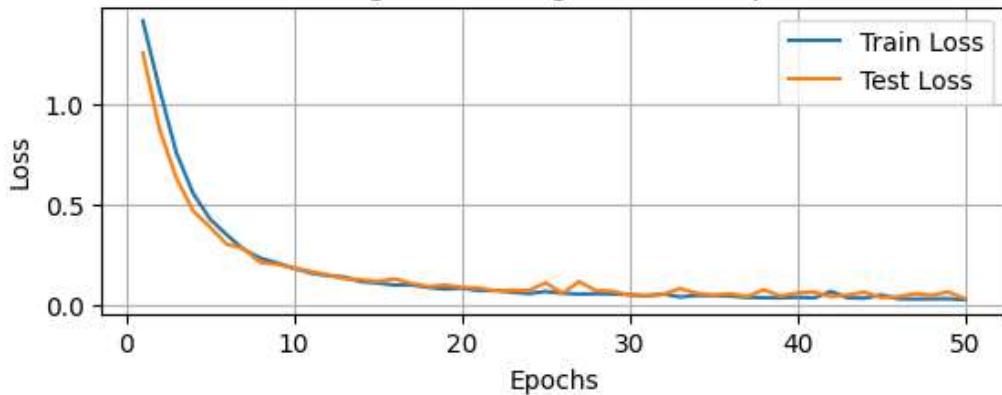
Train loss: 0.032123, Accuracy: 99.0%
Test Error:
Avg loss: 0.066229, Accuracy: 98.0%

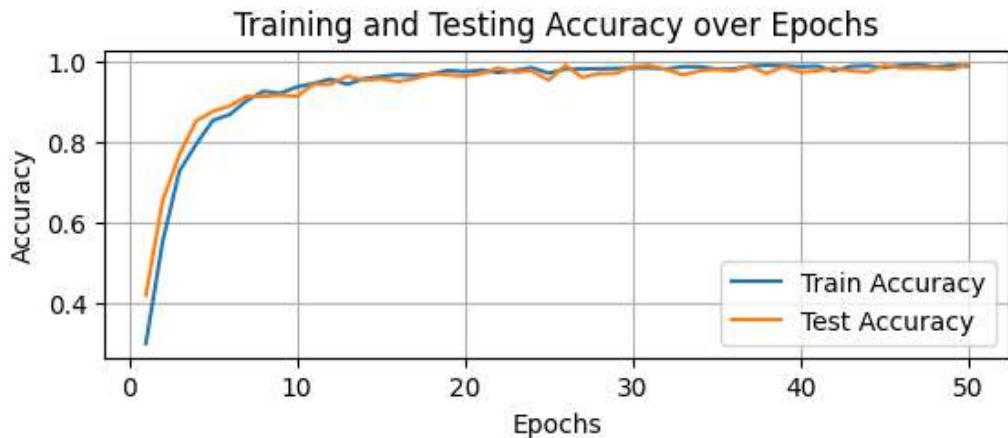
Epoch 50

Train loss: 0.027077, Accuracy: 98.8%
Test Error:
Avg loss: 0.031959, Accuracy: 99.3%

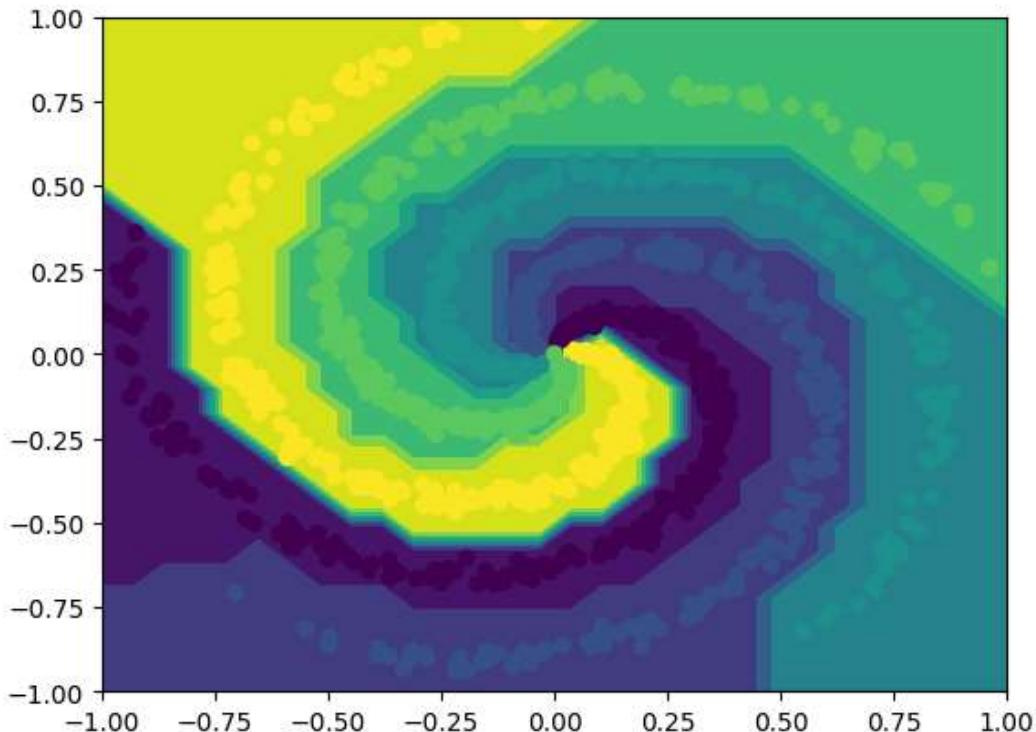
Done!

Training and Testing Loss over Epochs





```
In [8]: x,y = np.meshgrid(np.linspace(-1,1,30),np.linspace(-1,1,30))
xy = np.concatenate((x.reshape(-1,1),y.reshape(-1,1)),axis=1)
z = model(torch.tensor(xy).float()).detach().numpy()
z = np.argmax(z,1).reshape(30,30)
plt.contourf(x,y,z)
plt.scatter(train_data[:,0],train_data[:,1],c=train_labels)
plt.show()
```



2D feature space network

The network I have designed consists of two layers. Where the 2-dimension input goes to 256 neurons to 64 to the final 5 neurons, that is the output layer. I have drawn a lot of inspiration from exercises from week 3, which have provided a very stable network.

I have trained my network until it has converged. I have defined this as the current loss is higher than or equal to the average of the last 30 epochs. When the network has converged, I wait until the network gets the best performance on the training data.

The method I have chosen works really well for this scenario, but might not work optimally with different data. If there is a "big" local minima, the algorithm will think it has converged, and will

break when the accuracy will peak. This method works best with the ADAM optimizer, that uses momentum to push forward.

The network reaches up to 99.7% accuracy on the test data, which is quite good. The meshgrid plot also shows that the different classes are labeled correctly.